

Name	UB Number
Vaishnavi Varanasi	50590668

Project Update - MS Practicum MGS 649 SEM

Project Title: UB Research Assistant

Introduction:

The **UB Research Assistant** is an AI-powered academic tool designed specifically for the University at Buffalo community. It simplifies the research process for students, faculty, and scholars by offering quick paper summarization, smart recommendations, and a collaboration matching system with UB faculty. By combining the power of language models, semantic search, and open-access APIs, the tool helps users discover, digest, and connect around cutting-edge academic work.

Objectives:

- Automate the **summarization of research papers** using LLMs.
- Provide **real-time recommendations** of similar papers.
- Enable users to **find UB faculty members** working on related research areas.
- Reduce time spent on reading and searching through disconnected platforms.
- Lay the groundwork for future global collaboration beyond UB.

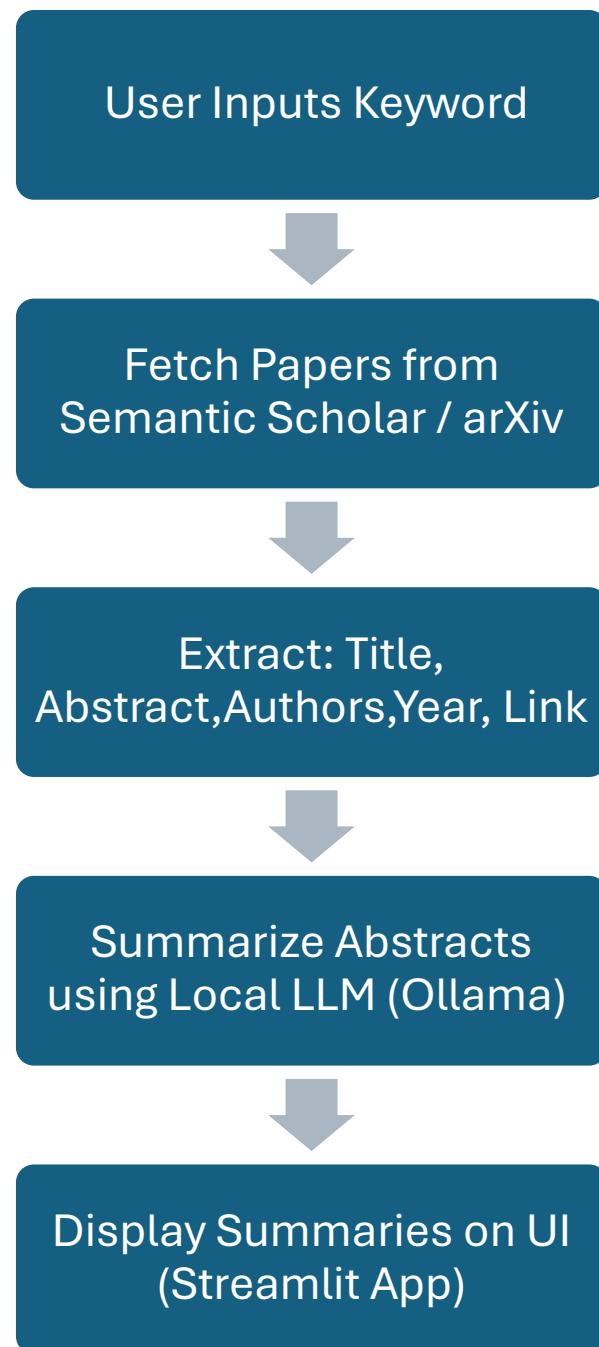
Scope:

1. **Research Paper Search**
 - Via Semantic Scholar and arXiv APIs.
2. **Paper Summarization**
 - Using a locally served language model (Mistral via Ollama).
3. **Paper Recommendation Engine**
 - Based on keyword relevance and embeddings.

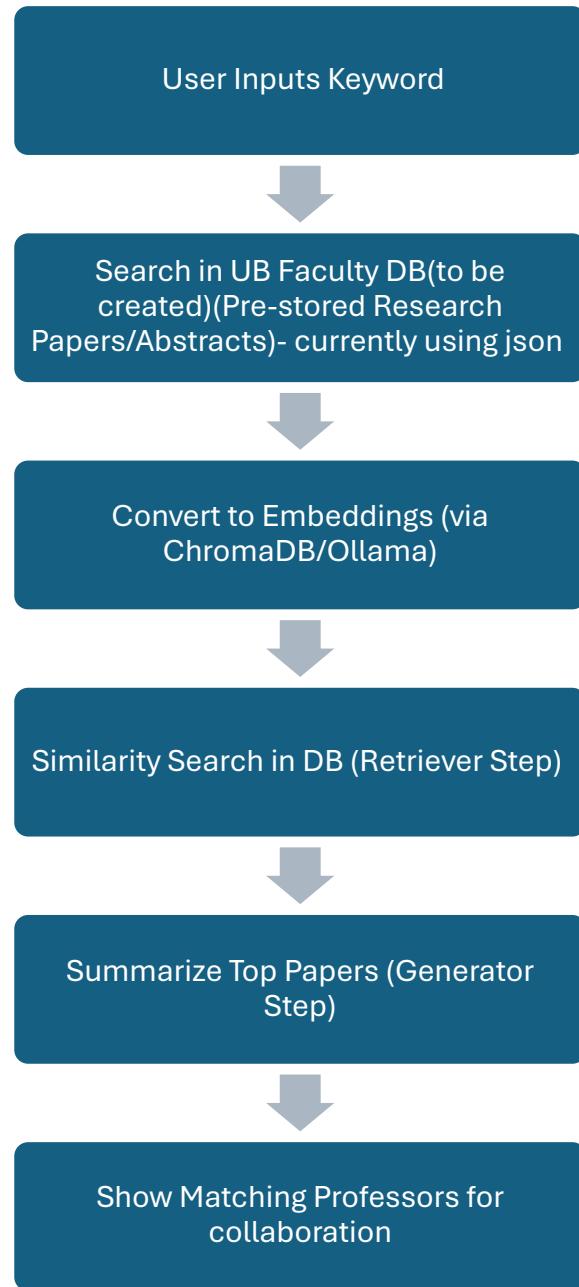
4. UB Faculty Research Matcher

- Matches user keywords with UB faculty research stored in a local JSON DB and ChromaDB.

Part A- LIVE SEARCH & ABSTRACT SUMMARIZATION



Part B- UB PROFESSORS' DATABASE & RAG COLLABORATION MATCHER



End-User Guide:

Step 1	Visit Streamlit App URL
Step 2	Enter Research keyword/topic
Step 3	In Part A- Fetch Papers Live -> Summarize and Display
Step 4	In Part B -> Search UB Faculty DataBase, Retrieve related Papers, Summarize, Show matching Professors

Technical Details of the Project:

1. Python- Version 3.11- Needed for compatibility with modern libraries like ChromaDB, Streamlit, and Ollama
2. Libraries Installed and their purpose-

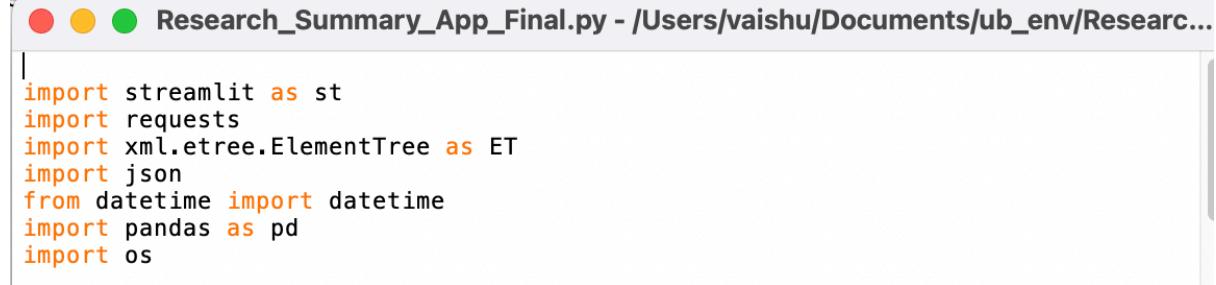
Library	Purpose
Streamlit	Frontend UI for easy, fast app deployment.
Chromadb	Vector database for embedding based retrieval
Requests	Fetch papers from external APIs
Ollama	Local LLM runtime without internet

3. APIs used-

API	Usage
Semantic Scholar	Highly relevant academic papers, structured metadata, free access for research
arXiv	Fetches open-access research papers across science, tech, and mathematics.

4. UI Framework choice(Streamlit)- It is extremely fast to develop and deploy a full-fledged web app with minimal Python code. It also auto-generates UI components (text boxes, buttons, tables) without needing HTML/CSS/JS knowledge. It is perfect for MVPs and internal academic tools.

Code:



```
import streamlit as st
import requests
import xml.etree.ElementTree as ET
import json
from datetime import datetime
import pandas as pd
import os
```

These imports load necessary Python libraries:

- streamlit: Builds the web app UI.
- requests: Makes API calls to fetch paper data.
- xml.etree.ElementTree: Parses XML from arXiv.
- json: Loads/saves JSON files.
- datetime: Timestamps feedback entries.
- pandas: For storing/handling feedback data.
- os: File path checking.

```
# ----- CONFIG -----
OLLAMA_MODEL = "mistral"
OLLAMA_URL = "http://localhost:11434/api/generate"
S2_URL = "https://api.semanticscholar.org/graph/v1/paper/search"
ARXIV_URL = "http://export.arxiv.org/api/query"
```

Defines constants:

- The model name for Ollama (LLM).
- Local API endpoint for generating summaries.
- Semantic Scholar and arXiv endpoints for paper data.

```
# ----- PAGE CONFIG -----
st.set_page_config(page_title="UB Research Assistant", layout="wide")
```

Sets the browser tab title and ensures full-width layout.

```
# ----- HEADER -----
st.markdown("""
    <div class="ub-header">
        <h1>UB Research Assistant</h1>
        <p>The UB way to research smarter</p>
    </div>
""", unsafe_allow_html=True)
```

Displays the blue header with project title and tagline.

```
# ----- SIDEBAR -----
with st.sidebar:
    st.markdown("""
        <div class="sidebar-content">
            <div class="sidebar-header">About This Assistant</div>
            <p>
                This tool acts as a personalized research assistant to students,
                <ul>
                    <li>Summarize recent papers</li>
                    <li>Explore UB faculty research</li>
                    <li>Find collaboration opportunities</li>
                </ul>
            </p>
            <div class="sidebar-header">Quick Links</div>
            <p>
                <a href="https://www.buffalo.edu/" target="_blank">UB Homepage</a>
                <a href="https://library.buffalo.edu/" target="_blank">UB Library</a>
                <a href="https://csed.buffalo.edu/" target="_blank">Dept. of Computer Science</a>
            </p>
            <div class="sidebar-header">Feedback</div>
            <p>We'd love to hear from you!</p>
        </div>
""", unsafe_allow_html=True)

with st.form("feedback_form"):
    name = st.text_input("Please enter your name")
    email = st.text_input("Email")
    feedback_text = st.text_area("Your Feedback", height=100)
    submitted = st.form_submit_button("Submit")

    if submitted and name.strip() and feedback_text.strip() and email.strip():
        timestamp = datetime.now().strftime("%Y-%m-%d %H:%M:%S")
        new_entry = pd.DataFrame([[timestamp, name.strip(), email.strip(), feedback_text.strip()]],
                                 columns=["Timestamp", "Name", "Email", "Feedback"])
        file_path = "feedback_log.csv"
        if os.path.exists(file_path):
            existing = pd.read_csv(file_path)
            updated = pd.concat([existing, new_entry], ignore_index=True)
        else:
```

In: 1 Col: 0

Sidebar shows:

- What the assistant does
- Quick links to UB resources
- A feedback form (name, email, message)

```
# ----- LOAD UB PAPERS -----
@st.cache_data
def load_ub_papers():
    with open("ub_papers.json", "r") as f:
        return json.load(f)

# ----- OLLAMA SUMMARIZATION -----
• Loads UB research papers from a local JSON file.
• Uses Streamlit caching for efficiency.
```

```
# ----- OLLAMA SUMMARIZATION -----
def summarize_text(text):
    prompt = f"Summarize the following research abstract in 2 simple sentences:\n{text}"
    response = requests.post(OLLAMA_URL, json={
        "model": OLLAMA_MODEL,
        "prompt": prompt,
        "stream": False
    })
    return response.json().get("response", "(No summary returned)").strip()

# ----- OLLAMA RELEVANCE EXPLANATION -----
• Sends abstract text to Ollama LLM API.
• Returns a 2-sentence summary.
```

```
# ----- OLLAMA RELEVANCE EXPLANATION -----
def get_relevance_reason(paper, keyword):
    prompt = (
        f"Given the following research paper and a keyword, explain in 1 line why\n"
        f"it relates to that keyword.\n\n"
        f"Keyword: {keyword}\n\n"
        f"Title: {paper['title']}\n"
        f"Abstract: {paper['abstract']}\n\n"
        f"Respond with a single sentence only."
    )
    try:
        response = requests.post(OLLAMA_URL, json={
            "model": OLLAMA_MODEL,
            "prompt": prompt,
            "stream": False
        })
        return response.json().get("response", "(No response)").strip()
    except Exception as e:
        return f"(Ollama error: {str(e)})"
    .....
```

- Sends paper details and user keyword to LLM.
- Returns 1-line explanation of why the paper is relevant.

```

# ----- SEMANTIC SCHOLAR FETCH -----
def fetch_semantic_scholar(query):
    res = requests.get(S2_URL, params={
        "query": query,
        "fields": "title,abstract,url,authors,year",
        "limit": 5
    })
    return [
        {
            "title": p["title"],
            "abstract": p.get("abstract", ""),
            "url": p.get("url", ""),
            "authors": ", ".join([author["name"] for author in p.get("authors", [])]),
            "year": p.get("year", "N/A")
        } for p in res.json().get("data", []) if p.get("abstract")
    ]
# ----- ARXIV FFTCH -----

```

- Sends query to Semantic Scholar.
- Returns top 5 papers with title, abstract, authors, year, and link.

```

# ----- ARXIV FETCH -----
def fetch_arxiv(query):
    url = f"{ARXIV_URL}?search_query=all:{query}&start=0&max_results=3"
    res = requests.get(url)
    root = ET.fromstring(res.content)
    ns = {'atom': 'http://www.w3.org/2005/Atom'}
    papers = []
    for entry in root.findall('atom:entry', ns):
        authors = [author.find('atom:name', ns).text for author in entry.findall(
            'atom:author', ns)]
        published = entry.find('atom:published', ns).text if entry.find('atom:published') else "N/A"
        year = published[:4] if published != "N/A" else "N/A"
        papers.append({
            "title": entry.find('atom:title', ns).text.strip(),
            "abstract": entry.find('atom:summary', ns).text.strip(),
            "url": entry.find('atom:id', ns).text.strip(),
            "authors": ", ".join(authors),
            "year": year
        })
    return papers

```

- Sends query to arXiv.
- Parses XML to extract papers with title, abstract, author names, year, and link.

```

# ----- MAIN TABS -----
summary_tab, ub_tab, collab_tab = st.tabs(["Summary", "UB Database", "Collaborat

```

Creates three tabs:

- Summary: Paper summarizer
- UB Database: Search UB-specific research

- Collaboration Finder: Find professors doing similar work

```
# ----- TAB 1: Summary -----
with summary_tab:
    st.markdown("### Research Summarizer")
    topic = st.text_input("Enter a keyword to fetch and summarize papers")

    if topic:
        st.info("Fetching papers from Semantic Scholar and arXiv...")
        papers = fetch_semantic_scholar(topic) + fetch_arxiv(topic)

        for paper in papers:
            summary = summarize_text(paper['abstract'])
            st.markdown(f"""
                <div class="paper-container">
                    <div class="paper-title">{paper['title']}</div>
                    <div class="paper-meta">
                        Authors: {paper['authors']}<br>
                        Year: {paper['year']}<br>
                        <a href="{paper['url']}'" target="_blank">View Full Paper</a>
                    </div>
                    <details><summary>Abstract</summary>
                    <p class="paper-abstract">{paper['abstract']}</p>
                    </details>
                    <div class="relevance-box"><b>AI Summary:</b> {summary}</div>
                </div>
            """, unsafe_allow_html=True)
```

- Lets user input a keyword
- Fetches papers from both APIs
- Uses LLM to summarize
- Displays output as paper cards with links and AI summaries

```

# ----- TAB 2: UB Paper Database -----
with ub_tab:
    st.markdown("### UB Research Paper Archive")
    st.markdown("Search and view relevant UB faculty research papers.")

keyword = st.text_input("Enter a research keyword", key="ub_search")

if keyword:
    ub_papers = load_ub_papers()
    filtered = []

    with st.spinner("Finding relevant UB papers..."):
        for paper in ub_papers:
            if keyword.lower() in paper['title'].lower() or keyword.lower():
                explanation = get_relevance_reason(paper, keyword)
                paper["reason"] = explanation
                filtered.append(paper)

    if filtered:
        st.markdown(f"**{len(filtered)} relevant papers found:**")
        for paper in filtered:
            st.markdown(f"""
                <div class="paper-container">
                    <div class="paper-title">{paper["title"]}</div>
                    <div class="paper-meta">
                        Authors: {', '.join(paper['authors'])}<br>
                        Department: {paper['department']}<br>
                        Email: {paper['email']}<br>
                        Year: {paper['year']}<br>
            """)

```

Ln: 1 Col: 0

- Searches the UB research database for matches to a keyword
- Shows paper details and an explanation of relevance

```

# ----- TAB 3: Collaboration Finder -----
with collab_tab:
    st.markdown("### UB Professor Collaboration Finder")
    st.markdown("Search for UB professors doing similar research.")

keyword = st.text_input("Enter a research keyword", key="collab_input")

if keyword:
    st.session_state["last_keyword"] = keyword
    ub_papers = load_ub_papers()
    results = [p for p in ub_papers if keyword.lower() in p["title"].lower()]
    if results:
        st.markdown(f"**{len(results)} matches found:**")
        for paper in results:
            explanation = get_relevance_reason(paper, keyword)
            st.markdown(f"""
                <div class="paper-container">
                    <div class="paper-title">{paper["title"]}</div>
                    <div class="paper-meta">
                        Authors: {', '.join(paper['authors'])}<br>
                        Department: {paper['department']}<br>
                        Email: {paper['email']}<br>
                        <a href="{paper['link']}' target="_blank">View Full
                    </div>
                    <div class="relevance-box"><b>Why Relevant:</b> {explanation}
                </div>
            """, unsafe_allow_html=True)
    else:
        st.warning("No matching faculty found.")
else:
    st.info("Enter a keyword to begin matching.")

```

- Matches user input keyword with UB faculty research
- Suggests professors and displays their work

Made it a Portfolio Project:

1. GitHub Repository
 - a. Full source code (Streamlit + Python)
 - b. CSS customization
 - c. API integration (Semantic Scholar + arXiv)
 - d. Sample UB faculty JSON database
2. Deployment Method

Local Deployment:

- a. Used streamlit run app.py to run locally
- b. Ollama runs locally for LLM summaries

Public URL:

- Deployed using Streamlit Cloud
- Steps:
 1. Pushed code to GitHub
 2. Connected to Streamlit Cloud
 3. Set up secrets (API key)
 4. Shared app via public URL: <https://research-assistant-at-ub.streamlit.app>

Challenges:

- Selecting the most flexible and open API (Semantic Scholar vs. arXiv)
- Handling abstract-less papers
- Designing a clean yet informative UI
- Limiting public access due to OpenAI's API rate limits

Screenshots

Set1:

Testing Semantic Scholar API:

```
# 📦 Import the requests library to make API calls
import requests
```

This line imports the requests library, which lets Python code talk to websites or APIs.

```
# 💡 Define a function that will search papers on Semantic Scholar
def fetch_from_semantic_scholar(query):
    print("\n🔍 SEMANTIC SCHOLAR RESULTS\n" + "-" * 40)

    # 🔎 Semantic Scholar API endpoint
    url = "https://api.semanticscholar.org/graph/v1/paper/search"

    # 🔍 Parameters for the search
    params = {
        "query": query, # 💬 Topic or keyword entered by the user
        "limit": 5, # 📄 Number of papers to fetch
        "fields": "title,abstract,authors,year,url" # 📄 What info we want back
    }
```

Enter your research topic: Applications of AI in Marketing Analytics

 SEMANTIC SCHOLAR RESULTS

1.  Title: Opportunities for the Use of AI in Marketing Communication by Educational Institutions
Abstract: Artificial intelligence has become an integral tool in various sectors, including marketing and marketing communication, offering unprecedented opportunities for enhancing efficiency and creativity. In marketing, AI's capabilities extend from automating routine tasks to providing deep insights through...
Authors: ['Ladislav Pátký', 'Jana Galera Matúšová', 'Katarína Načniaková']
Year: 2024
URL: <https://www.semanticscholar.org/paper/d3c8f3ef5dac7b8564bf08eb2edb02c24a82dd34>
2.  Title: Artificial Intelligence in Marketing Communication: A Comprehensive Exploration of the Integration and Impact of AI
Abstract: This study investigates the transformative impact of artificial intelligence (AI) on marketing communications through an evaluation research approach. Focused on enhancing personalization, efficiency, and strategic insight, the study explores AI's applications in content creation, customer service, ...
Authors: ['Hafize Nurgül Durmuş Şenyapar']
Year: 2024
URL: <https://www.semanticscholar.org/paper/0789e46115612436e37526dc7c6f8f644b576f69>
3.  Title: From analytics to empathy and creativity: Charting the AI revolution in marketing practice and education
Abstract: The rapid advancement of artificial intelligence (AI) increasingly demands an understanding of its impact on marketing practice and education. Our hybrid literature review synthesized 312 peer-reviewed articles on AI in marketing and consumer behavior, using scientometrics and the TCCM (Theory, Cont...
Authors: ['Giulia Pavone', 'Lars Meyer-Waarden', 'Andreas Munzel']
Year: 2024
URL: <https://www.semanticscholar.org/paper/43b901138fe4994ed58d700e7661aa426f32568a>
4.  Title: The Integration of AI in Digital Marketing: Opportunities and Challenges
Abstract: The integration of Artificial Intelligence (AI) in digital marketing has revolutionized the way businesses interact with customers and optimize their marketing strategies. AI technologies, such as machine learning, natural language processing, and predictive analytics, have enabled marketers to deli...
Authors: ['Tejaswini Bastry']
Year: 2025
URL: <https://www.semanticscholar.org/paper/aa852efdaaf07cf181e47479f0da9%c6edf628ba>
5.  Title: The impact of AI-supported marketing capabilities and analytics on SMEs' customer agility and marketing performance
Abstract: This research examines the impact of marketing analytics and artificial intelligence applications on customer agility and marketing performance in businesses that adopt e-commerce. In this quantitative study, data were collected through a questionnaire. Data collected from 227 managers online were a...
Authors: ['Fatma Demirag']
Year: 2025
URL: <https://www.semanticscholar.org/paper/37decead7ce84963b9a094b6dc8fea57af0d66b8>

(rag-env) (base) VaishnavisMBP2:rag-env vaishu\$ █

Set2:

```
# Insert documents into ChromaDB
collection.add(
    documents=documents,
    metadatas=metadatas,
    ids=ids
)

print("✅ Abstracts stored in ChromaDB")
```

Set3:

Live Summary Generator In Terminal:

```
import requests
import xml.etree.ElementTree as ET

OLLAMA_URL = "http://localhost:11434/api/generate"
OLLAMA_MODEL = "mistral"
QUERY = input("🔍 Enter a research topic:\n> ")
MAX_RESULTS = 3

# ◆ Step 1: Semantic Scholar
def fetch_semantic_scholar(query):
    print("Fetching from Semantic Scholar...")
    url = "https://api.semanticscholar.org/graph/v1/paper/search"
    params = {
        "query": query,
        "fields": "title,abstract,url",
        "limit": MAX_RESULTS
    }
    try:
        res = requests.get(url, params=params).json()
        return [
            {
                "source": "Semantic Scholar",
                "title": p["title"],
                "abstract": p.get("abstract", ""),
                "url": p.get("url", "")
            } for p in res.get("data", []) if p.get("abstract")
        ]
    except:
        return []

# ◆ Step 2: arXiv
def fetch_arxiv(query):
    print("Fetching from arXiv...")
    url = "https://export.arxiv.org/api/query?search_query=all:{query}&start=0&max_results={MAX_RESULTS}"
    res = requests.get(url)
    root = ET.fromstring(res.content)
    ns = {"atom": 'http://www.w3.org/2005/Atom'}
    papers = []
    for entry in root.findall('atom:entry', ns):
        title = entry.find('atom:title', ns).text.strip()
        summary = entry.find('atom:summary', ns).text.strip()
        link = entry.find('atom:id', ns).text.strip()
        papers.append({"source": "arXiv", "title": title, "abstract": summary, "url": link})
    return papers

# ◆ Step 3: CrossRef
def fetch_crossref(query):
    print("Fetching from CrossRef...")
    url = "https://api.crossref.org/works?query={query}&rows={MAX_RESULTS}"
    try:
        res = requests.get(url).json()
        return [
            {
                "source": "CrossRef",
                "title": item["title"][0],
                "abstract": item.get("abstract", "No abstract support"),
                "url": item.get("URL", "")
            } for item in res["message"]["items"] if "title" in item
        ]
    except:
        return []

# ◆ Step 4: Summarize with Ollama
def summarize_abstract(title, abstract, url, source):
    # Source : arXiv
    # Title : Bringing order into the realm of Transformer-based language models for
    # artificial intelligence and law
    # Summary : 1. Transformer-based Language Models (TLMs) are leading advancements in deep learning solutions, particularly in natural language processing and understanding. They have significantly improved AI approaches in various legal tasks.
    # 2. This article provides a comprehensive overview of TLM-based methods applied to AI problems in the legal field. It aims to understand how Transformers contribute to AI support in legal processes, while also identifying current limitations and opportunities for further research.
    # Link : http://arxiv.org/abs/2308.05502v2

    # Skipping CrossRef paper 'Artificial Intelligence in Trauma' (no abstract).
    # Skipping CrossRef paper 'NIHR Research: Artificial Intelligence' (no abstract).
    # Skipping CrossRef paper 'Intelligence, Artificial' (no abstract).

    # (rag-fresh-env) (base) VaishnavisMBP2:rag-fresh-env vaishu$ Computer Networks
    # bash: Computer: command not found
    # (rag-fresh-env) (base) VaishnavisMBP2:rag-fresh-env vaishu$ python live_summary_multiple_APIs.py
    # Enter a research topic:
    # > Computer Networks

    # Fetching from Semantic Scholar...
    # Fetching from arXiv...
    # Fetching from CrossRef...

    # Summarizing 6 papers using Ollama...

    # Source : arXiv
    # Title : Computing Power Network: A Survey
    # Summary : 1. A new networking paradigm called Computing Power Network (CPN) is proposed to leverage distributed computing resources in the era of cloud computing, edge computing, and smart devices. This network aims to connect heterogeneous resources for flexible scheduling.
    # 2. This research survey covers the state-of-the-art on CPN, discussing its architecture, advantages, issues such as modeling, resource allocation, network forwarding, and more. It also presents a testbed realization, use cases, enabling technologies, and open challenges for future research.
    # Link : http://arxiv.org/abs/2210.06080v2

    # Source : arXiv
    # Title : Analysis of network by generalized mutual entropies
    # Summary : 1. A new method called generalized mutual entropy is proposed to analyze complex network structures, which has been tested on computer-simulated scale-free networks, random networks, and their combinations.
    # 2. Potential real-world applications of this method in network analysis are explored.
    # Link : http://arxiv.org/abs/0709.0929v1

    # Source : arXiv
    # Title : A Multilayer Model of Computer Networks
    # Summary : 1. This research proposes a formal model of computer networks using hierarchical multilayer networks, where each layer is represented as a multiplex network.
    # 2. The study also identifies the essential set of layers needed for network architecture representation when considering the system methodology in network analysis.
    # Link : http://arxiv.org/abs/1509.00721v1

    # Skipping CrossRef paper 'Computer networks, the journal, and computer networks, the technology' (no abstract).
    # Skipping CrossRef paper 'Computer Networks: Publisher's Note' (no abstract).
    # Skipping CrossRef paper 'Local computer networks' (no abstract).

    # (rag-fresh-env) (base) VaishnavisMBP2:rag-fresh-env vaishu$
```

In Streamlit:

The screenshot shows the Streamlit interface for the UB Research Assistant. On the left, a sidebar contains sections for "About This Assistant", "Quick Links" (with links to UB Homepage, UB Libraries, and Dept. of Computer Science), and "Feedback". A feedback form with a placeholder "Please enter your name" is also present. The main content area has a blue header bar with the title "UB Research Assistant" and the tagline "The UB way to research smarter". Below the header, there are three tabs: "Summary" (which is active and highlighted in red), "UB Database", and "Collaboration Finder". The "Research Summarizer" section contains a search bar with the placeholder "Enter a keyword to fetch and summarize papers".

The screenshot shows the Streamlit interface for the UB Research Assistant, specifically the "UB Database" tab. The layout is identical to the "Summary" tab, with the sidebar on the left and the main content area on the right. In the main content area, the "UB Research Paper Archive" section is displayed. It includes a search bar with the placeholder "Enter a research keyword" containing the text "blockchain". Below the search bar, it says "1 relevant papers found:" followed by a card for the paper "Blockchain for Public Health Records". The card displays the authors (Dr. Vaishnavi Varanasi), department (Health Informatics), email (vvaranasi@buffalo.edu), and year (2022). A "View Full Paper" link is provided. A note states: "This paper proposes a blockchain-based solution to enhance the security and accessibility of electronic health records for underserved populations." A callout box highlights the relevance of the paper to the keyword "blockchain".

The screenshot shows a web application titled "The UB way to research smarter". At the top right are "Deploy" and three-dot buttons. Below the title is a navigation bar with "Summary", "UB Database", and "Collaboration Finder" (which is underlined). The main content area is titled "UB Professor Collaboration Finder" and contains a search bar with the placeholder "Enter a research keyword" and the input "blockchain". A message "1 matches found:" is followed by a card for "Blockchain for Public Health Records". The card includes author information ("Authors: Dr. Vaishnavi Varanasi", "Department: Health Informatics", "Email: vvaranasi@buffalo.edu"), a "View Full Paper" link, and a summary: "Why Relevant: The "Blockchain for Public Health Records" research is relevant to the keyword "blockchain" as it presents a practical application of blockchain technology in managing secure and accessible electronic health records."

About This Assistant

This tool acts as a personalized research assistant to students, faculty and researchers

- Summarize recent papers
- Explore UB faculty research
- Find collaboration opportunities

Quick Links

[UB Homepage](#)
[UB Libraries](#)
[Dept. of Computer Science](#)

Feedback

We'd love to hear from you!

Please enter your name

The screenshot shows a web application titled "The UB way to research smarter". At the top right are "Stop" and "Deploy" buttons. Below the title is a navigation bar with "Summary" (underlined), "UB Database", and "Collaboration Finder". The main content area is titled "Research Summarizer" and contains a search bar with the placeholder "Enter a keyword to fetch and summarize papers" and the input "healthcare Ai". A message "Fetching papers from Semantic Scholar and arXiv..." is displayed.

About This Assistant

This tool acts as a personalized research assistant to students, faculty and researchers

- Summarize recent papers
- Explore UB faculty research
- Find collaboration opportunities

Quick Links

[UB Homepage](#)
[UB Libraries](#)
[Dept. of Computer Science](#)

Feedback

We'd love to hear from you!

Please enter your name

The screenshot shows the main interface of the UB Research Assistant. At the top right, there are buttons for 'RUNNING...', 'Stop', 'Deploy', and a three-dot menu. Below the header, a blue banner reads 'UB Research Assistant' and 'The UB way to research smarter'. A navigation bar with tabs 'Summary' (which is red), 'UB Database', and 'Collaboration Finder' is visible. The main content area is titled 'Research Summarizer'. It features a search bar with the placeholder 'Enter a keyword to fetch and summarize papers' and a text input field containing 'healthcare Ai'. Below the search bar, a message says 'Fetching papers from Semantic Scholar and arXiv...'. A summary card for a paper is displayed, titled 'Multimodal Healthcare AI: Identifying and Designing Clinically Relevant Vision-Language Applications for Radiology'. The card includes author information: Nur Yildirim, Hannah Richardson, M. Wetscherek, Junaid Bajwa, Joseph Jacob, M. Pinnock, Stephen Harris, Daniel Coelho De Castro, Shruthi Bannur, Stephanie L. Hyland, Pratik Ghosh, M. Ranjit, Kenza Bouzid, Anton Schwaighofer, Fernando P'erez-Garc'ia, Harshita Sharma, O. Oktay, M. Lungren, Javier Alvarez-Valle, A. Nori, Anja Thieme. The year is listed as 2024. There are buttons for 'View Full Paper' and 'View Paper Summary'. On the left side of the interface, there is a sidebar with sections for 'About This Assistant', 'Quick Links' (with links to UB Homepage, UB Libraries, and Dept. of Computer Science), and 'Feedback' (with a text input field for name). A message at the bottom of the sidebar says 'We'd love to hear from you!'. A text input field for 'Please enter your name' is also present.

This screenshot shows a detailed view of a paper summary within the UB Research Assistant interface. At the top right, there are buttons for 'RUNNING...', 'Stop', 'Deploy', and a three-dot menu. The main content area displays a summary of a paper by Stephen Harris, Daniel Coelho De Castro, Shruthi Bannur, Stephanie L. Hyland, Pratik Ghosh, M. Ranjit, Kenza Bouzid, Anton Schwaighofer, Fernando P'erez-Garc'ia, Harshita Sharma, O. Oktay, M. Lungren, Javier Alvarez-Valle, A. Nori, and Anja Thieme, published in 2024. The summary discusses the integration of AI and vision-language models for radiology, mentioning tasks like generating findings based on medical images and answering visual questions. It highlights the clinical utility of these capabilities and the design process used. Below the summary, there is a section for 'AI Summary' which provides a brief overview of the study's findings and their implications for healthcare AI. On the left side, there is a sidebar with sections for 'About This Assistant', 'Quick Links' (with links to UB Homepage, UB Libraries, and Dept. of Computer Science), and 'Feedback' (with a text input field for name). A text input field for 'Please enter your name' is also present.

References:

1. Semantic Scholar API Documentation

<https://api.semanticscholar.org/>

(Used for fetching academic paper metadata like title, abstract, authors, etc.)

2. arXiv API Documentation

<https://info.arxiv.org/help/api/index.html>

(Used for accessing open-access science and technology research papers)

3. Streamlit – Official Docs

<https://docs.streamlit.io/>

(Used to build the front-end interface of the app)

4. Ollama – Open Source LLM Serving

<https://ollama.com/>

(Used for running language models like Mistral locally for summarization tasks)

5. ChromaDB Documentation

<https://docs.trychroma.com/>

(Used as a vector database to store and retrieve research embeddings for matching)

6. RAG (Retrieval-Augmented Generation) Overview

Lewis, Patrick, et al. "Retrieval-augmented generation for knowledge-intensive NLP tasks."

NeurIPS 2020. <https://arxiv.org/abs/2005.11401>

(The foundational method your system uses for improved paper recommendations)